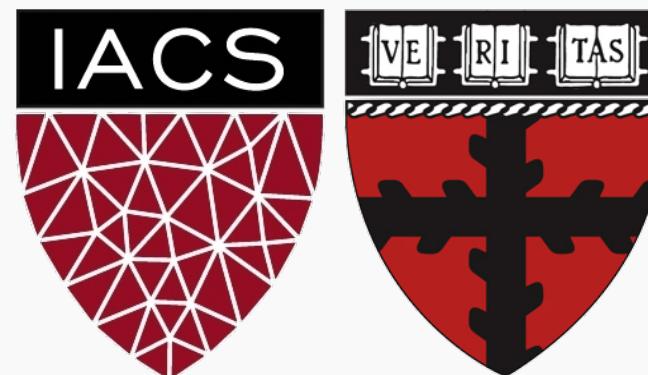


Lecture 22: Interpretation of Prediction Models

CS109A Introduction to Data Science
Pavlos Protopapas, Kevin Rader and Chris Tanner



ANNOUNCEMENTS

- Homework 7 OH:
 - For conceptual questions: Faculty will continue their office hours.
 - If you have problems with TensorFlow please let us know on ED. We will arrange special OH to help if necessary.
- Project:
 - Milestone3 due today. EDA and base model. Talk to your assigned TF!
- AC 209A
 - Advanced Section today.



Outline

Variable Importance

Interpretation through Predictions

Adding Uncertainty



Variable Importance



Variable Importance for Tree-Based Models

How does sklearn determine **variable importance** (`feature_importance`) from a tree-based model?

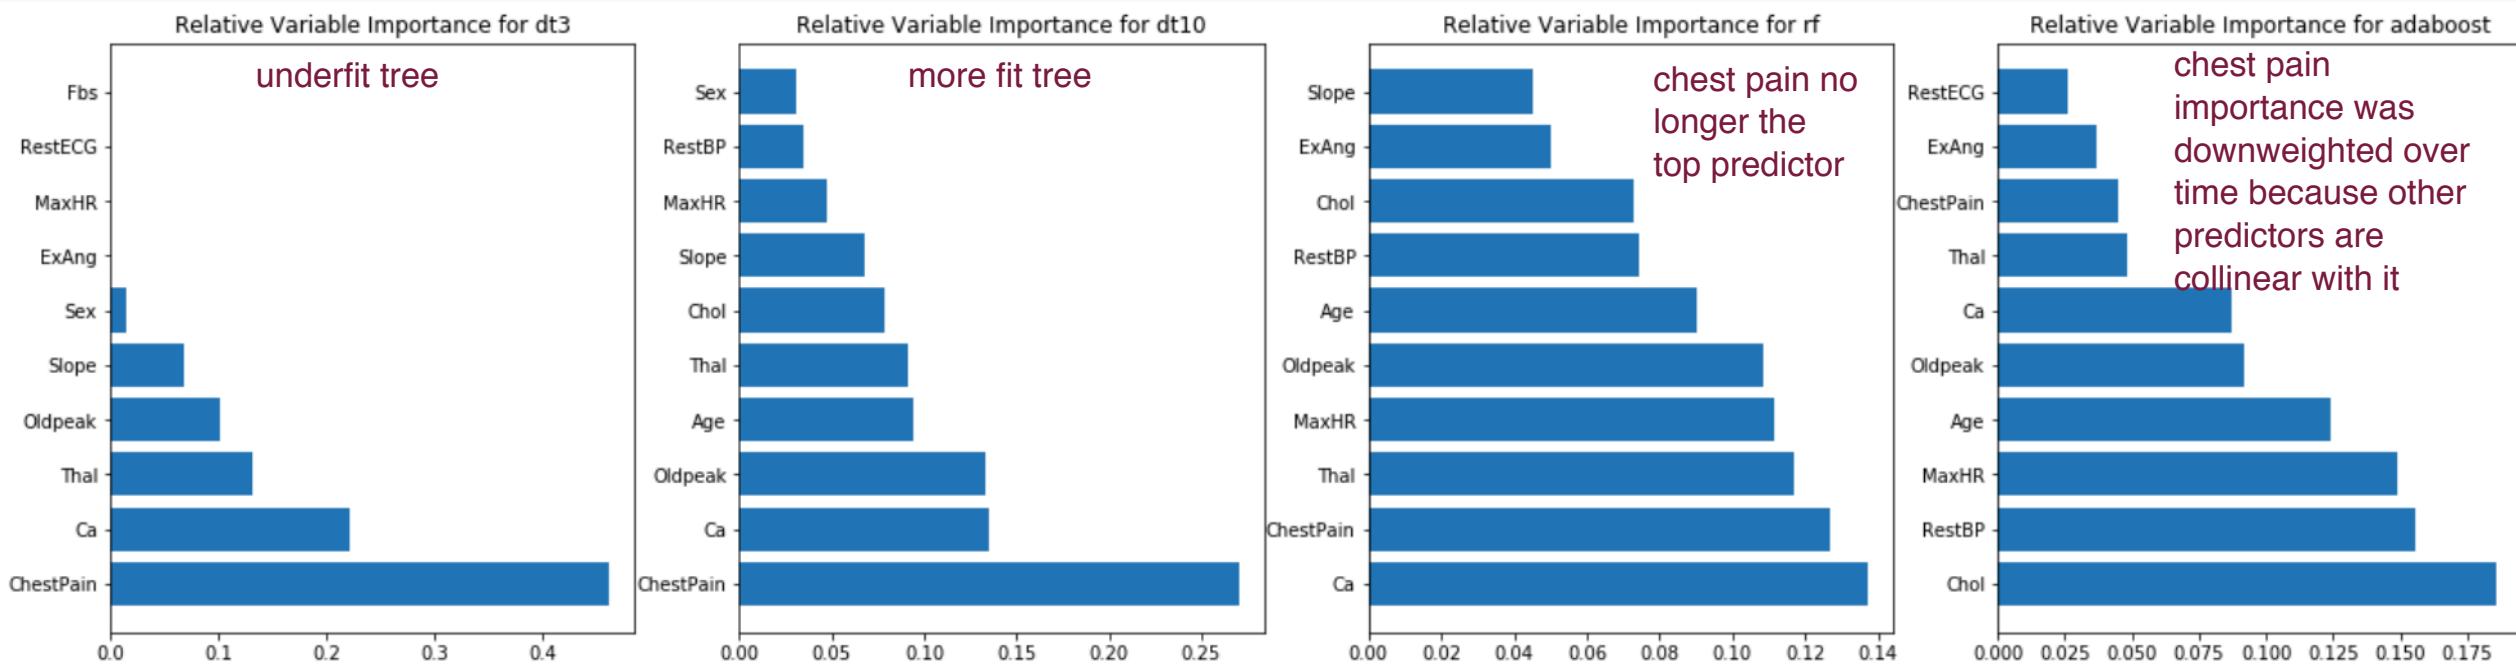
- It determines the improvement in the loss function every time a predictor is involved in a split.
- More specifically, it calculate the total amount that the SSE (for regression) or Gini index (for classification) is improved (decreased) due to splits over a given predictor (averaged over all B trees if a bagged/random forest method).

How should variable importance compare across the various different tree models we've considered (trees, random forests/bagging, and boosting)?

A picture is worth a thousand words...

Variable Importance for trees, bags, and boosts

Below are the variable importance plots for the top 10 predictors for each of a (i) decision tree with maxdepth=3, (ii) decision tree with maxdepth=10, (iii) a random forest, and (iv) an adaboost classifier.



Compare them? Are the differences surprising?

Feature Importance in a Neural Network

How can one measure feature importance in a Neural Network?

shuffling data to permute relationship between
predictor and output - will neural net still do well
without this predictor?

Its not so easy 😞. What could potentially be done?

We can calculate the predictions from a neural net, and then fit a decision tree model to the predicted values.

What can go wrong with this approach?

variable importance measurements are ‘variable’ between the tree approaches

Other Variable Importance Measures

What other approaches can be taken to measure variable importance?

Alternative:

- Record the prediction accuracy on the *oob* samples for each tree.
- Randomly permute the data for column j in the *oob* samples and record the accuracy again.
- The decrease in accuracy as a result of this permuting is averaged over all trees, and is used as a measure of the importance of variable j in the random forest.

Permutation Variable Importance

This idea of re-permuting a variable and refitting a model to see how much more poorly it performs is called **permutation feature importance**.

It is sometimes preferred to the standard feature importance, why? standardized

When two features are correlated and one of the features is permuted, the model will still have access to the feature through its correlated feature.

What is the one glaring disadvantage to this?

Computational time, and things may not ‘add up’ SSEs might not all add up at the end unlike for variable importance...rescale?

Permutation Variable Importance in sklearn

To perform permutation variable importance in sklearn (or keras), one can use the ELI5 library:

<https://eli5.readthedocs.io/en/latest/index.html>

This is easy to do with sklearn models, but not easy to do with tensorflow (but it can be done).

An example is worth a thousand words:



Permutation Variable Importance Example

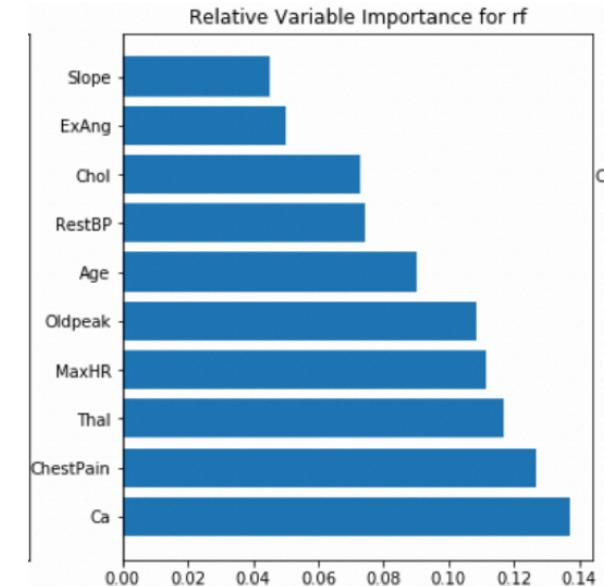
you don't have to refit the model if you permute the variable - just look at the worsening in predictions with the variable permuted

```
#permutation importance
from eli5.sklearn import PermutationImportance
from eli5.permutation_importance import get_score_importances

perm = PermutationImportance(randomforest).fit(X_test, y_test)
#eli5.show_weights(perm, feature_names=X.columns)
print(X.columns)
eli5.explain_weights(perm)
```

```
Index(['Age', 'Sex', 'ChestPain', 'RestBP', 'Chol', 'Fbs', 'RestECG', 'MaxHR',
       'ExAng', 'Oldpeak', 'Slope', 'Ca', 'Thal'],
      dtype='object')
```

Weight	Feature
0.1082 ± 0.0161	x12
0.0689 ± 0.0435	x7
0.0393 ± 0.0675	x11
0.0361 ± 0.0636	x2
0.0328 ± 0.0359	x8
0.0295 ± 0.0245	x9
0.0295 ± 0.0131	x1
0.0197 ± 0.0245	x3 chest pain is collinear with other predictors
0.0164 ± 0.0688	x10
0.0131 ± 0.0131	x6
0.0098 ± 0.0161	x4
0 ± 0.0000	x5
-0.0098 ± 0.0334	x0



The problem with Variable Importance

Variable Importance is great! It tells you what features are important in shaping the model.

But what is missing?

- It does not give any measure for how the predictors are related to the response (positive, negative, quasi-linear, curved, interactions, etc.).
- This is where the parametric model wins out! Inference and interpretations are much easier and the whole point in these models.

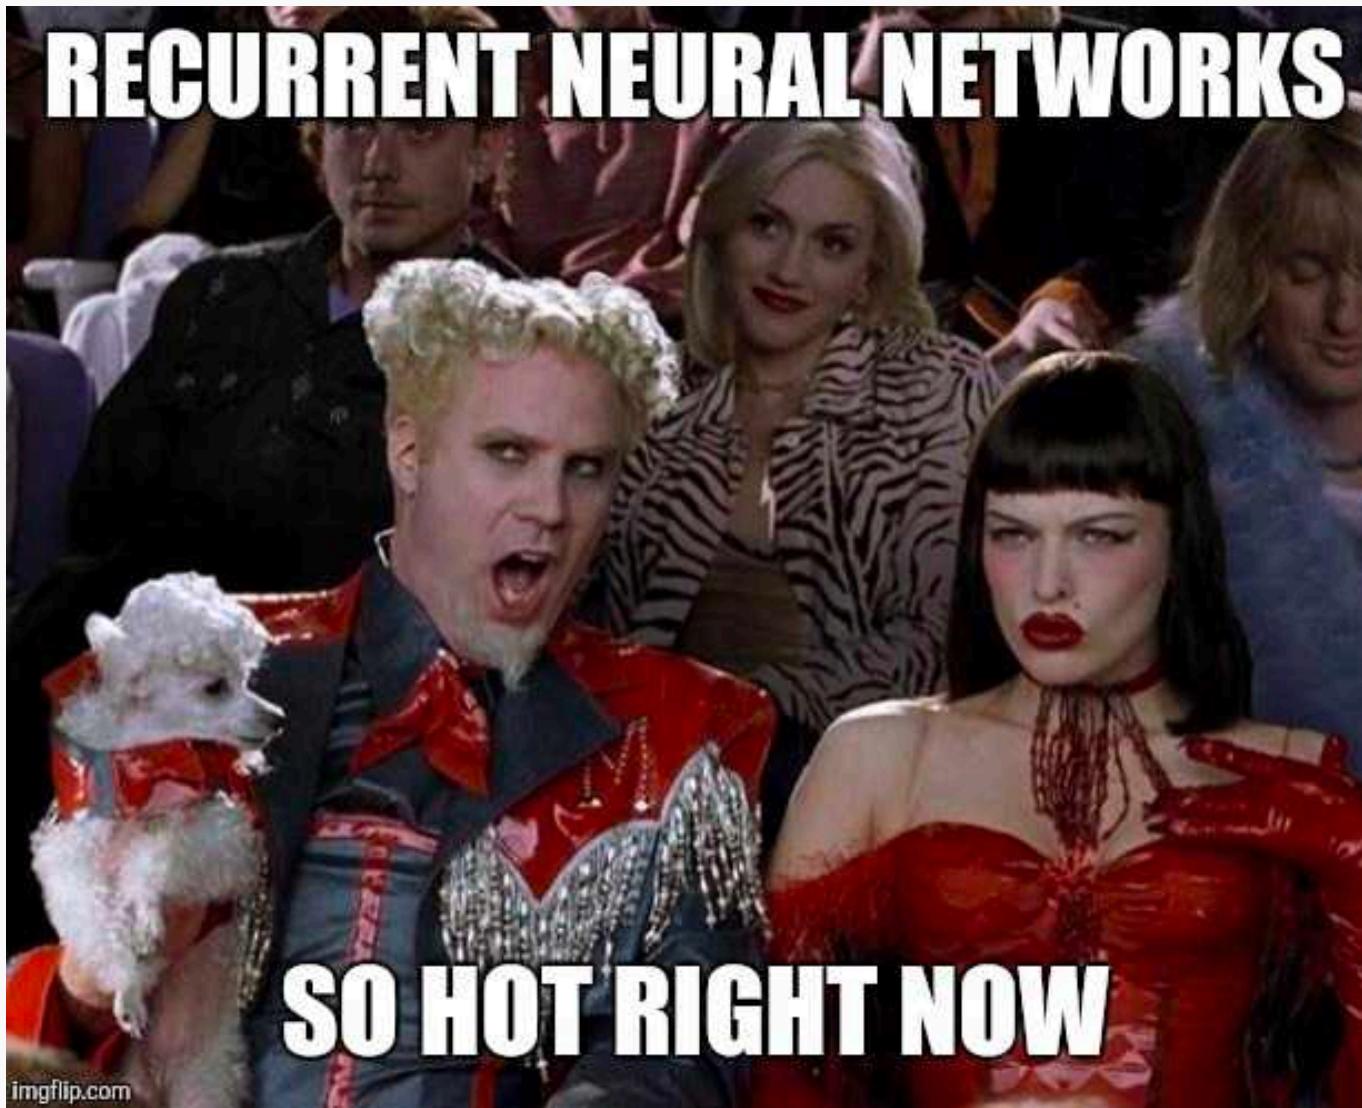
What can we do to measure these relationships in a machine learning or non-parametric model? Think: what did we do with k -NN?

hone in on effect of one variable, holding the other variables constant - we can use visuals for this

What needs to be done algorithmically to put this in practice?

RECURRENT NEURAL NETWORKS

SO HOT RIGHT NOW



imgflip.com



Interpretation through Predictions



Parametric vs. Nonparametric models

In a machine learning model (like ensemble methods and neural networks), the association between predictors and the response are not measured directly as these models are ‘black box’ models:

Inputs (X , predictors). → black box (NN, sklearn, etc.) → Outputs (Y , response)

What if we care about how the predictors relate to the response? This is where we need to figure out what the black box is doing to transform the inputs into the outputs.

Simplest Approach: observed \hat{Y} vs. X_j

Use predict (or better yet, predict_proba) to plot the observed predicted values vs. the observed values for X_j .

What is a problem with this approach how can we fix it?

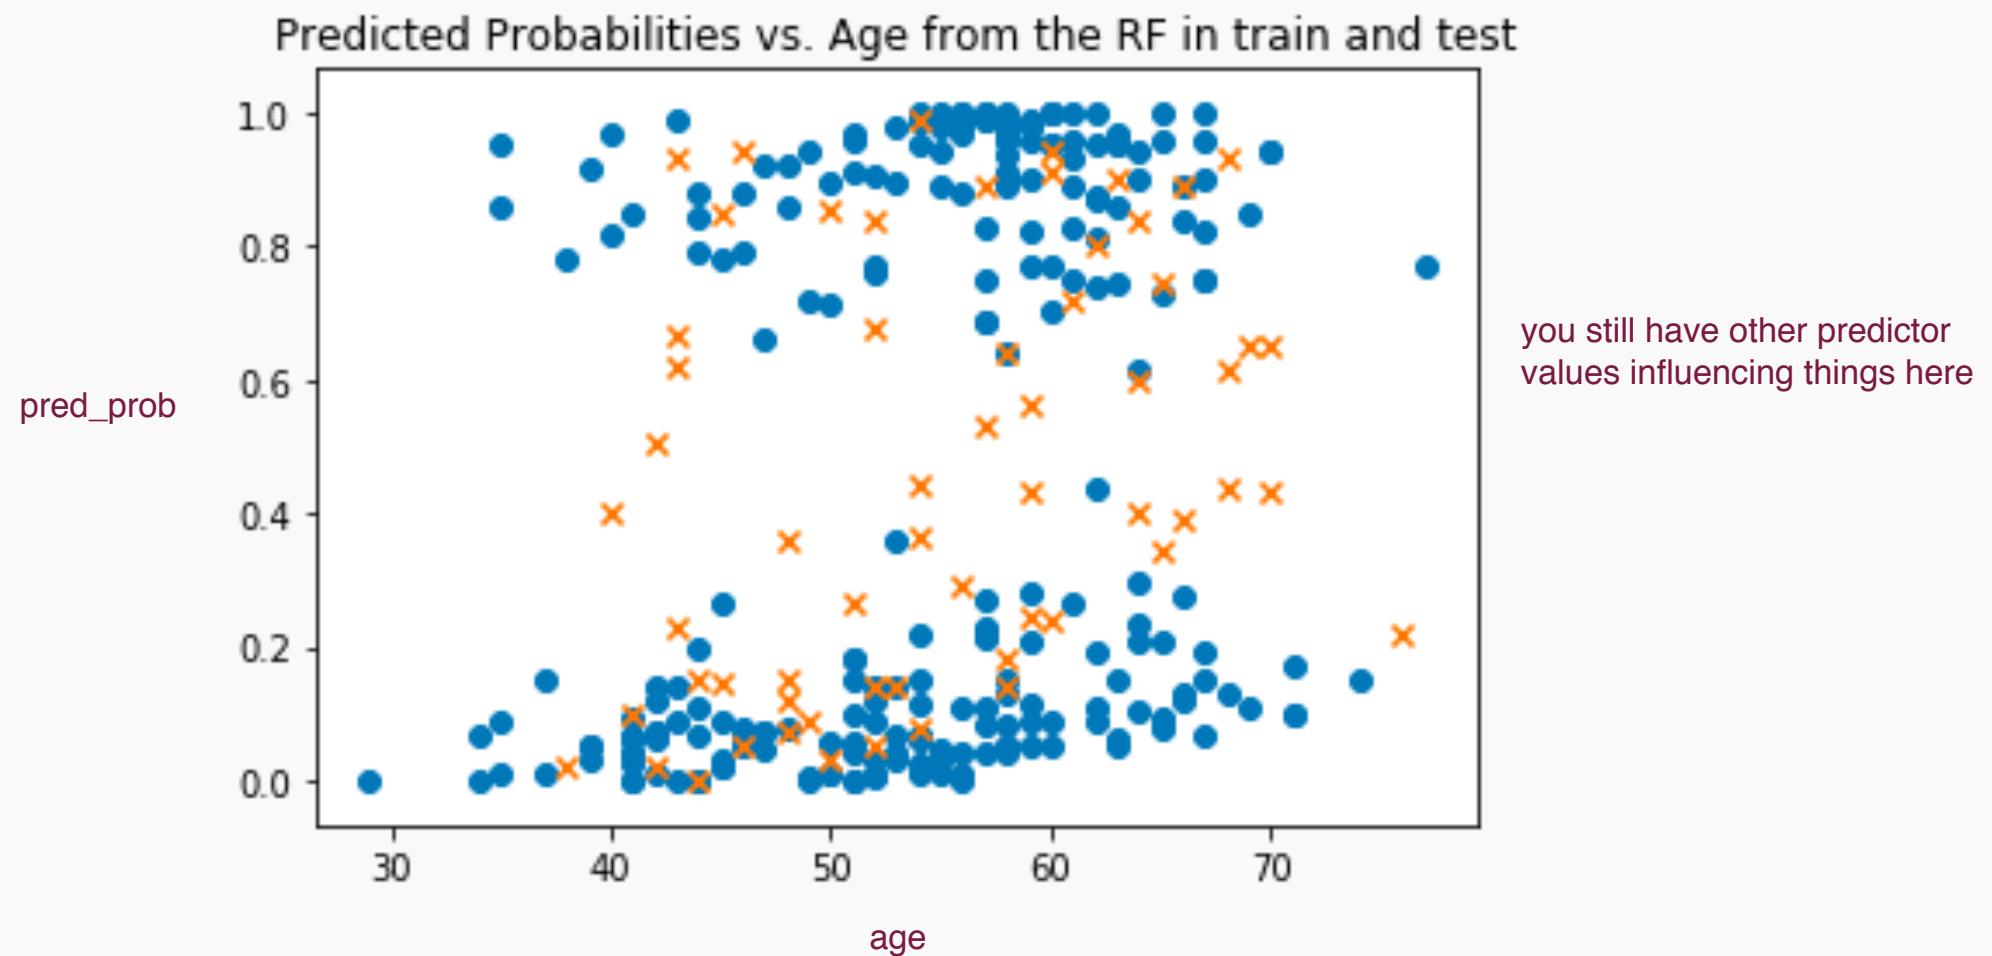
not holding the other variables constant

The fix is not so easy. We cannot just fit a logistic regression model so easily to the predicted probabilities. Why not?

An example is worth a thousand words...



Example: observed \hat{Y} vs. X_j



Unboxing the black box

Inputs (X , predictors). → black box (NN, sklearn, etc.) → Outputs (Y , response)

This gives us our approach: vary the inputs (the predictors) and see what happens to the response.

If we care about the ‘marginal’ or ‘conditional’ effect of how a specific X relates to Y , then we should vary only one predictor at a time.

same idea for coefficients in multiple regression model

How should we handle the other predictors? That is to say, what value should we keep them at? common aggregate value?



Unboxing the black box (cont.)

There are two general approaches to interpreting the machine learning model through predictions:

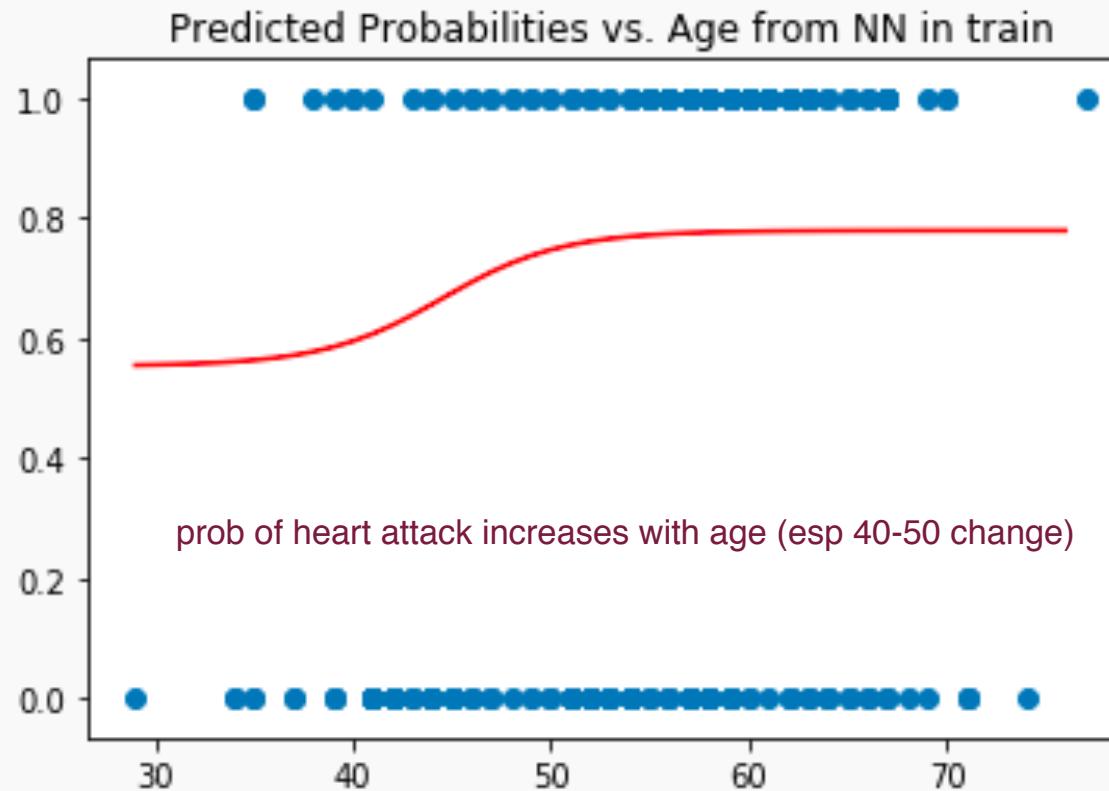
The approach is just like in multiple regression: what is the marginal effect of a unit change in X_j holding all the other predictors constant.

So at what values should we hold the other predictors?

There are two general approaches **holding the other predictors constant**:

1. Predict \hat{Y} at the **mean (or most common) value** for each of the other predictors, vary only the predictor you care about, X_j , and plot the predictions \hat{Y} vs. X_j . this eliminates interactive effects
2. Predict \hat{Y} at the **observed values of** for all the other predictors, **vary only the predictor you care about, X_j** , and plot the predictions \hat{Y} vs. X_j . Essentially this means creating a new data frame for each observation, and imputing all reasonable values of X_j in.

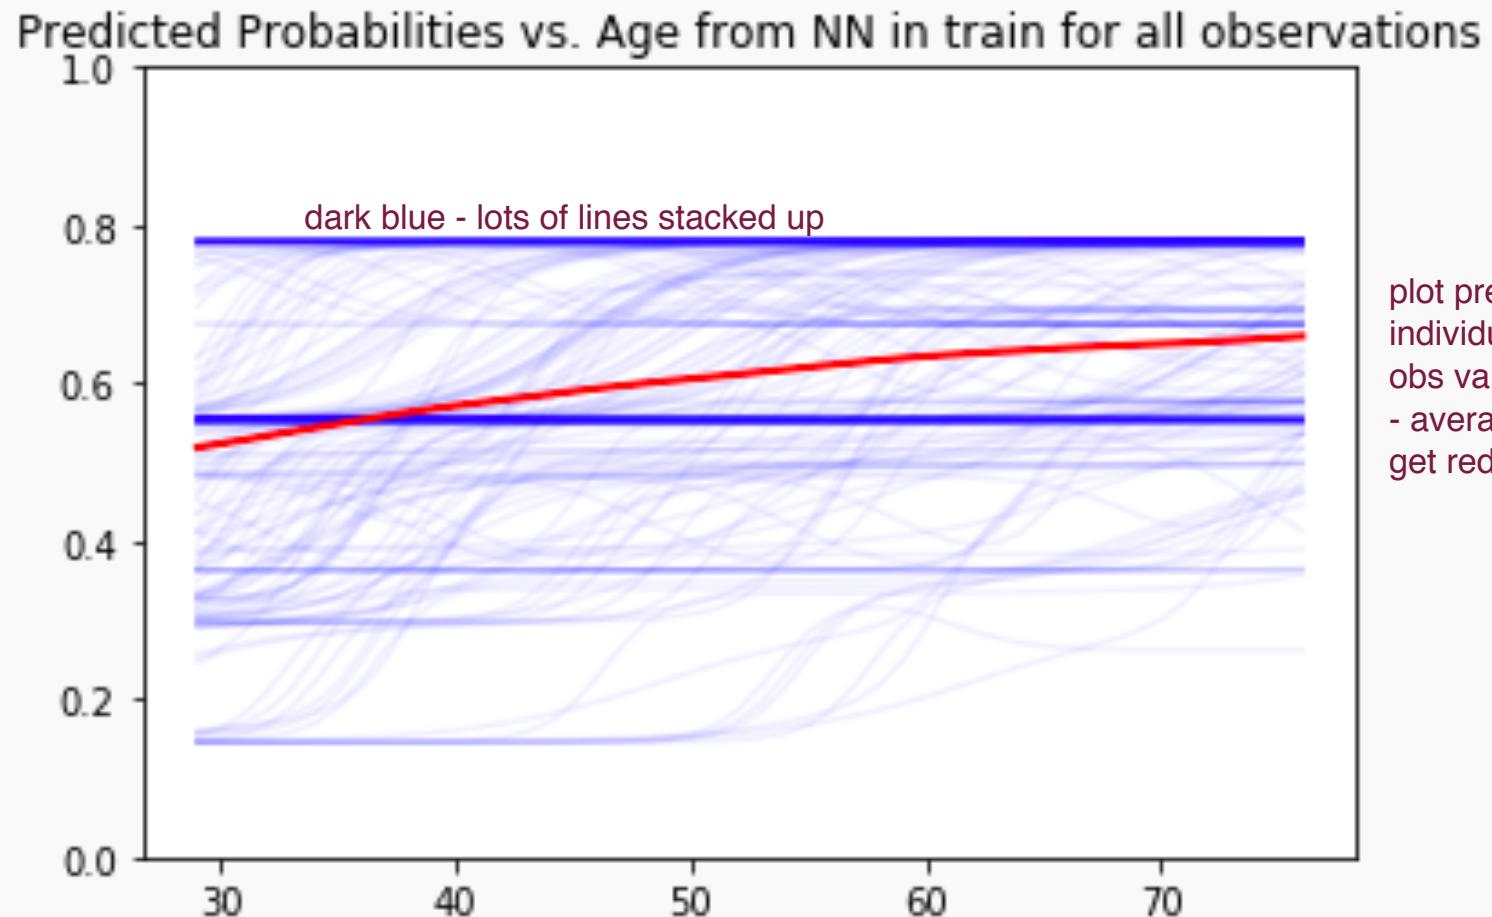
Predicted probabilities vs. Age at the means



Interpret this plot. What does this say for how Cardiac Arrest is associated with Age?

this prediction may not apply to all individuals - this is just for one individual, with the other predictors held at some aggregate value (not observed values)

Predicted probabilities vs. Age for all observations



Interpret this plot. What does this say for how Cardiac Arrest is associated with Age?

What if we want the joint relationship with two predictors?

How can we look at how the response relates to two predictors at once?

2D map of colors?

This is a bit trickier to do. Why?

We can go back to our friend: the classification boundary!

What if we want the joint relationship with two predictors?



Should I use PCA components?

What if we want an overall picture of how the response relates to the predictors?

Well, this is where PCA components sometimes come in to play.

useful for clustering? evaluating how strong your model is?

But this is no Bueno for interpretability. Why? linear combinations of observed values

Adding Uncertainty



Adding Uncertainty

Not only should we plot predictions, but we can (should) also include the uncertainty of these predictions.

put SEs around our estimate - how do we get SE?

For example, there may be a big difference between.

a model that is certain that $P(y=1) = 0.5$ and one that is less certain

How can we add uncertainty measures to these predictions? How do we do it for linear/logistic regression?

Bootstrapping

parametric models - have certain formulas for SEs
non-parametric models - bootstrap

BOOTSTRAPPING

I DONT THINK IT MEANS WHAT YOU
THINK IT MEANS

memegenerator.net



Bootstrapping Basics, review

How do we perform a bootstrap? What are the key steps?

The bootstrap algorithm is an approach to mimic the uncertainty of taking the observed sample from a population. So we:

1. Sample the same number of observations as was observed (in train).
2. Sample with replacement (so we get some randomness).

Why does this work?

our observed sample is the best guess as to what our population looks like
we are trying to mimic resampling from the population



Bootstrapping Basics, review

When is a bootstrapping approach used?

1. When a probabilistic closed form solution is not known or not easy to get at.
2. When the assumptions for a parametric model break down.
3. When there is no ‘formula’ at all to add the uncertainty into a calculation.
4. Or we are just being lazy.

Bootstrapping Basics, review

How do we use the results of a bootstrap technique?

It provides us a re-created **estimate** every time we bootstrap.

This estimate could be a mean, a beta, a plot of average predictions, etc.

But they will almost always be for a summary, not a single observation, and thus the sample size is already taken care of (the n in the formula). Aka, do not divide by $\text{sqrt}(n)$ again!

uncertainty already incorporated into our bootstrap sampling

Bootstrapping for Predictions

So we can refit a model on a bootstrap sample of data, and look at the predicted probabilities each time.

Then we can take the top 97.5th percentile and bottom 2.5th percentile to build a 95% confidence interval for the predictions!

But how does this work for random forests and bagging models?

Yikes!

bootstrap is already baked in to these models